




PrismLearning.AI

Not a chatbot. Lumi drives your document itself.

 **Built with AMD Instinct GPUs x Fireworks AI**

Passive reading is dead.



Cognitive Load

Flat PDFs create a dense wall of unabsorbable text, quickly exhausting student retention capacities and focus spans.



Passive Dumping

Standard AI chatbots simply hand over answers without guidance, short-circuiting critical thinking and conceptual learning.



Hallucination Risks

Unconstrained LLMs easily pull students completely out of the core curriculum by generating unsupported factual errors.

Lumi Drives the Document — Not the Chat

Not a chatbot. Not a flashcard app.

- ✓ **Grounded RAG:** A pristine pipeline strictly locked within the uploaded document bounds to maintain 100% curriculum bounds.
- ✓ **Scaffolded UX:** The AI refuses to give final answers, strategically probing and guiding the student step-by-step to answers.
- ✓ **Agentic Control:** Every reply is JSON that drives the UI — scroll, highlight, unlock, gamify. Never just chat.

The screenshot displays the Lumi AI interface, which is designed for educational use. On the left, a sidebar titled 'Workspace' contains sections for 'Documents' (with a '+ Add' button), 'Curriculum' (listing topics like Problem, Solution, Architecture & Infrastructure, etc.), 'Assessments' (with buttons for 'Generate Flashcards' and 'Generate Quiz'), and 'Settings'. At the bottom of the sidebar is an 'Export PDF' button. The main content area is divided into two parts. The top part, titled '1. Problem', discusses 'Key challenges with traditional study tools' and includes a diagram showing how 'Static PDF', 'Standard Chatbot', and 'Unconstrained LLM' lead to 'Cognitive Load', 'Passive Answers', and 'Hallucinations' respectively, all of which result in 'Low Retention'. The bottom part, titled '2. Solution', describes the 'PrismLearning.AI' experience through three pillars: 'Grounded RAG', 'Scaffolded UX', and 'Agentic Control'. On the right, a chat window shows a conversation with 'Lumi, Your AI tutor'. The chat includes a 'Step 4 of 5' indicator, a 'Grounded RAG explanation' section, and a 'we want to avoid passive answers?' prompt. The chat also features a 'Yes' button, a 'Ofcourse!' button, and a 'Nope what does rag mean?' button. The chat window ends with an input field for 'Answer Lumi or ask a question...'.

Back to Dashboard

Workspace

Documents (1) + Add

PrismLearning AI 13 Slide Pl... Learn

Curriculum

1. Problem

2. Solution

3. Architecture & Infrastructure

4. Core Features & Learning Engine

5. Safety & Trust

6. Market, Pricing & Roadmap

Assessments

Flashcard Deck (0)

Generate Flashcards

Generate Quiz

Settings

Export PDF

Complete a concept with Lumi to earn flashcards to export.

Reading level Academic

1. Problem

Key challenges with traditional study tools

- **Cognitive Load** – Flat PDFs present a dense wall of text, quickly exhausting retention and focus.
- **Passive Dumping** – Standard AI chatbots hand over answers without guidance, bypassing critical thinking.
- **Hallucination Risks** – Unconstrained LLMs can generate unsupported factual errors, pulling learners off-curriculum.

Static PDF → Cognitive Load

Standard Chatbot → Passive Answers

Unconstrained LLM → Hallucinations

Cognitive Load → Low Retention

Passive Answers → Low Retention

Hallucinations → Low Retention

2. Solution

PrismLearning.AI delivers an **agentic, curriculum-grounded tutoring experience** through three pillars:

- **Grounded RAG** – Retrieval-augmented generation locked to the uploaded document, guaranteeing 100% curriculum bounds.
- **Scaffolded UX** – The AI refuses final answers, instead probing step-by-step to guide learners.
- **Agentic Control** – JSON payloads drive UI actions (scrolling, unblurring) for a live, interactive canvas.

Lumi
Your AI tutor

Spark - Lv 1

Step 4 of 5

Grounded RAG explanation

we want to avoid passive answers?

Yes

Great! The third issue is **hallucination risks** – unconstrained LLMs can invent facts that aren't in the curriculum, pulling you off-track. Imagine a tutor confidently stating a wrong date in history; it feels real but misleads you. Does that example illustrate why we need to guard against hallucinations?

Ofcourse!


Now we'll look at the first pillar of the solution: **Grounded RAG** – the system only pulls information from the uploaded document, so every answer stays within the curriculum. Think of it like a librarian who only hands you books that are actually on the shelf you requested, never inventing new titles. Does that analogy make sense to you?

Nope what does rag mean?

Sure! **RAG** stands for **Retrieval-Augmented Generation** – the system first pulls (retrieves) relevant text from your uploaded document, then uses that material to generate its answer, so it never invents information outside the curriculum. Does that clarify why the tutor stays "grounded" to the source?

Answer Lumi or ask a question...

Why We Win



The Engineering Stack



Frontend

Next.js 14, Tailwind CSS, and Framer Motion for sleek glassmorphic transitions.

01



State Engine

Zustand state engine catches JSON payloads for instant, low-latency DOM manipulation.

02



Backend

Python FastAPI & Pydantic designed for highly performant text extraction from documents.

03



Data & Auth

Supabase PostgreSQL with Clerk auth; uploads stay in memory only, never stored.

04

Powered by **Fireworks AI & AMD**

Traditional Cloud APIs

- ✗ High latency ruins real-time agentic interaction.
- ✗ Prohibitively expensive inference costs limit growth.
- ✗ Breaks immersion of a fluid, live conversational tutor.

Fireworks AI on AMD Instinct™

- ✓ Blazing fast Tokens-Per-Second output capacity.
- ✓ Sub-second UI mutation latency for live updates.
- ✓ Serverless scaling on AMD Instinct™ GPUs, low cost.

One architecture, three models — gpt-oss-120b for tutoring, Gemma 3 27B for flashcards, Gemma 4 for Enterprise, all served on AMD Instinct™ GPUs.

The Active Learning Canvas



Fog of War

The core text remains locked behind frosted glass, unblurring section-by-section as student mastery is proven.



Walkie-Talkie

Low-latency, vocal interaction driven securely by the Web Speech API for natural conversation.



Printable Decks

Instantly compiles smart micro-quizzes into cleanly formatted PDF flashcard cutouts for physical study.

Beyond Flashcards: The Learning Engine

Confidence

Before revealing a quiz answer, students self-rate their confidence — feedback then flags over- or under-confidence.

Boss Battles

Mastering a chapter unlocks a themed, timed practice exam that proves the material actually stuck.

Mastery Decay

Flashcards resurface on a spaced-repetition schedule; mastery quietly decays until reviewed again.

Built for Academic Integrity



Data Sovereignty

Zero-retention data framework. Uploaded curriculum assets are fully discarded post-session and never used for base training.



Anti-Cheat Guardrails

The Agentic Tutor is hard-prompted to never write essays, output source copy blocks, or hand over final homework answers directly.



Hallucination-Free

Strict structural system prompts constrain the model to instruct solely from parameters defined in your syllabus.

A Scalable SaaS Engine

MOST POPULAR

Free Tier

Basic exploration

Free

- ✓ Ingest standard PDFs
- ✓ Core active tutor chat

Pro Tier

Designed for power students

\$10/mo

- ✓ Advanced UI limits unlocked
- ✓ Unlimited syllabus uploads
- ✓ Full Walkie-Talkie voice tutor


Campus Tier

Institutional standard

B2B


- ✓ Dedicated Gemma 4 on AMD Instinct™
- ✓ Teacher Analytics dashboards

The Future of PrismLearning



Enough talking.

Let's step inside the Prism.



The Visionaries Behind PrismLearning

We are looking for seed-stage feedback and beta testers to break our Agentic Engine.



patrickilagan.dev@gmail.com

John Patrick Ilagan - Student Founder & Full-Stack Architect

*"Built out of personal necessity because
current study tools are fundamentally broken."*

